



Chapter 5: Advanced SQL

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan
See www.db-book.com for conditions on re-use



Outline

- Accessing SQL From a Programming Language
- Functions and Procedures
- Triggers
- ☀ Recursive Queries
- ☀ Advanced Aggregation Features



Accessing SQL from a Programming Language

A database programmer must have access to a general-purpose programming language for at least two reasons

- Not all queries can be expressed in SQL, since SQL does not provide the full expressive power of a general-purpose language.
- Non-declarative actions -- such as printing a report, interacting with a user, or sending the results of a query to a graphical user interface -- cannot be done from within SQL.



Accessing SQL from a Programming Language (Cont.)

There are two approaches to accessing SQL from a general-purpose programming language

- A general-purpose program -- can connect to and communicate with a database server using a collection of functions
- Embedded SQL -- provides a means by which a program can interact with a database server.
 - The SQL statements are translated at compile time into function calls.
 - At runtime, these function calls connect to the database using an API that provides dynamic SQL facilities.



ODBC

- Open DataBase Connectivity (ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - open a connection with a database,
 - send queries and updates,
 - get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC
- 可參考課本圖5.5的ODBC範例。
- 在Windows裡, 有“ODBC資料來源管理員”, 可看到電腦裡已經灌好的 driver library。



JDBC SUBSECTIONS

- Connecting to the Database
- Shipping SQL Statements to the Database System
- Exceptions and Resource Management
- Retrieving the Result of a Query
- Prepared Statements
- ※ Callable Statements
- ※ Metadata Features
- ※ Other Features
- ※ Database Access from Python



JDBC

- JDBC is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.
- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the statement object to send queries and fetch results
 - Exception mechanism to handle errors
- 以下的程式範例必須先import 「java.sql.*」 .



JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
    try (Connection conn = DriverManager.getConnection(
        "jdbc:oracle:thin:@db.yale.edu:2000:univdb", userid, passwd);
        Statement stmt = conn.createStatement();
    )
    {
        ... Do Actual Work ....
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle);
    }
}
```

**NOTE: Above syntax works with Java 7, and JDBC 4 onwards.
Resources opened in “try (...)” syntax (“try with resources”) are automatically closed at the end of the try block**



JDBC Code (Cont.)

- Update to database

```
try {  
    stmt.executeUpdate(  
        "insert into instructor values('77987', 'Kim', 'Physics', 98000)");  
} catch (SQLException sqle)  
{  
    System.out.println("Could not insert tuple. " + sqle);  
}
```

- Execute query and fetch and print results

```
ResultSet rset = stmt.executeQuery(  
    "select dept_name, avg (salary)  
     from instructor  
     group by dept_name");  
  
while (rset.next()) {  
    System.out.println(rset.getString("dept_name") + " " +  
                      rset.getFloat(2));  
}
```

dept_name	avg(salary)
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000



JDBC Code Details

- Getting result fields:
 - **rs.getString(“dept_name”)** and **rs.getString(1)** equivalent if **dept_name** is the first argument of select result.
- Dealing with Null values

```
int a = rs.getInt("a");
if (rs.wasNull()) System.out.println("Got null value");
```



Prepared Statement

- We can create a prepared statement in which some values are replaced by “?”, thereby specifying that actual values will be provided later.
- The database system compiles the query when it is **prepared**. Each time the query is **executed**, the database system reuses the compiled form of the query and applies the new values as parameters.
- Prepared statements allow for more **efficient** execution in cases where the same query can be compiled once and then run multiple times with different parameter values.
- Example:
 - **PreparedStatement pStmt = conn.prepareStatement("insert into instructor values(?,?,?,?,?)");**
pStmt.setString(1, "88877");
pStmt.setString(2, "Perry");
pStmt.setString(3, "Finance");
pStmt.setInt(4, 125000);
pStmt.executeUpdate();
pStmt.setString(1, "88878");
pStmt.executeUpdate();



Prepared Statement (Cont.)

- Example:
 - “`insert into instructor values(' " + ID + " ', ' " + name + " ', ' " + dept_name + " ', " + balance + ")`”
- What if name is “Perry”?
 - “`insert into instructor values(' " + “88877” + " ', ' " + “Perry” + " ', ' " + ...`
 - ⇒ `insert into instructor values('88877 ', ‘Perry’, ‘.....`
- What if name is “D’Souza”?
 - “`insert into instructor values(' " + “88877” + " ', ' " + “D’Souza” + " ', ' " + ...`
 - ⇒ `insert into instructor values('88877 ', ' D’Souza ', ‘.....`

單引號配對出錯!!
- WARNING: always use prepared statements when taking an input from the user and adding it to a query
 - 因為單引號會失去特殊功能 (看下一頁)
 - NEVER create a query by concatenating strings



SQL Injection

- Suppose query is constructed using
 - “**select * from instructor where name = “” + name + “”**
- Suppose the user, instead of entering a name, enters:
 - X' or 'Y' = 'Y
- then the resulting statement becomes:
 - "select * from instructor where name = “” + "X' or 'Y' = 'Y" + “”"
 - which is:
 - select * from instructor where name = ‘X’ or ‘Y’ = ‘Y’ => 看到全部
 - User could have even used
 - X‘; update instructor set salary = salary + 10000; -- => 更改資料
- Prepared statement internally uses:
“select * from instructor where name = 'X\' or \'Y\' = \'Y'
 - **Always use prepared statements, with user inputs as parameters**
- SQL injection 實例:
<https://news.ltn.com.tw/news/life/breakingnews/2971370>



補充PDO (PHP) prepared statements example

```
<?php
    //連接資料庫
    $db = new
    PDO('mysql:host=localhost;dbname=test;charset=utf8',$user,$password);
    //設定想要新增入資料庫的資料內容如下
        $dept = "資工";
        $no = "00557888";
        $name = "發發發";
        $club = "桌遊社";
    //設定要使用的SQL指令
        $query = ("insert into student values(?, ?, ?, ?)");
        $stmt = $db->prepare($query);
    //執行SQL語法
        $result = $stmt->execute(array($dept,$no,$name,$club));
?
>
```



補充 PDO – query example

```
<?php
    include "db_conn.php";
    $no = "00557888";
    $query = ("select * from student where no = ?");
    $stmt = $db->prepare($query);
    $error = $stmt->execute(array($no));
    $result = $stmt->fetchAll();
    for($i=0; $i<count($result); $i++){
        echo "dept:". $result[$i]['dept']. ' ';
        "no:". $result[$i]['no']. ' ';
        "name:". $result[$i]['name']. ' ';
        "club:". $result[$i]['club']. ' ';
        '<br>';
    }
}
```



Transaction Control in JDBC

- By default, each SQL statement is treated as a separate transaction that is committed automatically
 - bad idea for transactions with multiple updates
- Can turn off automatic commit on a connection
 - `conn.setAutoCommit(false);`
- Transactions must then be committed or rolled back explicitly
 - `conn.commit();` or
 - `conn.rollback();`
- `conn.setAutoCommit(true)` turns on automatic commit.



❖ Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, C++, Java, Fortran, and PL/1,
- A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded* SQL.
- The basic form of these languages follows that of the System R embedding of SQL into PL/1.
- **EXEC SQL** statement is used in the host language to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement>;

Note: this varies by language:

- In some languages, like COBOL, the semicolon is replaced with END-EXEC
- In Java embedding uses # SQL { };



Functions and Procedures



Functions and Procedures

- Functions and procedures allow “business logic” to be stored in the database and executed from SQL statements.
- These can be defined either by the procedural component of SQL or by an external programming language such as Java, C, or C++.
- The syntax we present here is defined by the SQL standard.
 - Most databases implement nonstandard versions of this syntax.



Declaring SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (dept_name varchar(20))
    returns integer
begin
    declare d_count integer;
    select count (*) into d_count
    from instructor
    where instructor.dept_name = dept_name;
    return d_count;
end
```

instructor (ID, name, dept_name, salary)
department (dept_name, building, budget)

- The function *dept_count* can be used to find the department names and budget of all departments with more than 12 instructors.

```
select dept_name, budget
from department
where dept_count (dept_name ) > 12;
```

- 函數可直接用在select和where裡。



※ Table Functions

- The SQL standard supports functions that can return tables as results; such functions are called **table functions**

- Example: Return all instructors in a given department

```
create function instructor_of (dept_name char(20))  
    returns table ( ID varchar(5),  
                    name varchar(20),  
                    dept_name varchar(20),  
                    salary numeric(8,2))  
  
return table  
(select ID, name, dept_name, salary  
     from instructor  
   where instructor.dept_name = instructor_of.dept_name);
```

- Usage

```
select *  
from table (instructor_of ('Music'));
```



SQL Procedures

- The *dept_count* function could instead be written as procedure:

```
create procedure dept_count_proc (in dept_name varchar(20),
                                   out d_count integer)
```

```
begin
```

```
    select count(*) into d_count
        from instructor
       where instructor.dept_name = dept_count_proc.dept_name
```

```
end
```

- The keywords **in** and **out** are parameters that are expected to have values assigned to them and parameters whose values are set in the procedure in order to return results.
- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
declare d_count integer;
call dept_count_proc( 'Physics', d_count);
```

- 補充: SQL Server和MySQL支援stored procedure.



Language Constructs for Procedures & Functions

- SQL supports constructs that gives it almost all the power of a general-purpose programming language.
 - Warning: most database systems implement their own variant of the standard syntax below.
- Compound statement: **begin ... end**,
 - May contain multiple SQL statements between **begin** and **end**.
 - Local variables can be declared within a compound statements
- While and repeat statements:
 - **while** boolean expression **do**
 sequence of statements ;
end while
 - **repeat**
 sequence of statements ;
until boolean expression
end repeat



※ Language Constructs (Cont.)

- **For loop**
 - Permits iteration over all results of a query
- Example: Find the budget of all departments

```
declare n integer default 0;
for r as
    select budget from department
do
    set n = n + r.budget
end for
```



Language Constructs – if-then-else

- Conditional statements (**if-then-else**)

```
if boolean expression
    then statement or compound statement
elseif boolean expression
    then statement or compound statement
else statement or compound statement
end if
```

- 複雜的例子請參考課本圖5.8
- ※ MariaDB function: <https://mariadb.com/kb/en/create-function/>
- ※ MariaDB procedure: <https://mariadb.com/kb/en/create-procedure/>



※ External Language Routines

- SQL allows us to define functions in a programming language such as Java, C#, C or C++.
 - Can be more efficient than functions defined in SQL, and computations that cannot be carried out in SQL can be executed by these functions.
- Declaring external language procedures and functions

```
create procedure dept_count_proc(in dept_name varchar(20),
                                out count integer)
```

language C

external name '/usr/avi/bin/dept_count_proc'

```
create function dept_count(dept_name varchar(20))
```

returns integer

language C

external name '/usr/avi/bin/dept_count'

- Benefit: efficiency; drawback: security



Triggers



Triggers (觸發程序)

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- To design a trigger mechanism, we must:
 - Specify the conditions under which the trigger is to be executed.
 - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
 - Syntax illustrated here may not work exactly on your database system; check the system manuals



Triggering Events and Actions in SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to a relation or a specific attribute; e.g.,
 - **after update on takes <- 較多軟體支援**
 - **after update of grade on takes**
- Values of attributes before and after an update can be referenced
 - **referencing old row as** : for deletes and updates
 - **referencing new row as** : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints. For example, convert blank grades to null.

```
create trigger setnull_trigger before update on takes
referencing new row as nrow
for each row
when (nrow.grade = ' ')
begin atomic
    set nrow.grade = null; } } declare a transaction
end;
```



Trigger Example

- E.g. `time_slot_id` is not a primary key of `timeslot`, so we cannot create a foreign key constraint from `section` to `timeslot`.
- Alternative: use triggers on `section` and `timeslot` to enforce integrity constraints

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
    select time_slot_id
    from time_slot) /* time_slot_id not present in time_slot */
begin
    rollback
end;
```

- 完整範例請參考課本圖5.9

```
section (course_id, sec_id, semester, year, building, room_number, time_slot_id)
time_slot (time_slot_id, day, start_time, end_time)
```



練習

- Write triggers to enforce the referential integrity constraint from *section* to *time_slot*, on **deletion** to *time_slot*.

(hint: You need to check that the *time_slot_id* of the tuple being deleted is either still present in *time_slot*, or that no tuple in *section* contains that particular *time_slot_id* value; otherwise, referential integrity would be violated.)

Answer:



Trigger to Maintain credits_earned value

- create trigger *credits_earned* after update of *grade* on *takes* referencing new row as *nrow* referencing old row as *orow* for each row
when (*nrow.grade* <> 'F' and *nrow.grade* is not null)
and (*orow.grade* = 'F' or *orow.grade* is null)
begin atomic
update *student*
set *tot_cred*= *tot_cred* +
(select *credits*
from *course*
where *course.course_id*= *nrow.course_id*)
where *student.id* = *nrow.id*;
end;

takes					
<i>student</i>	<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
	1	CS-101	1	Spring	2010
				F	A
<i>course</i>			<i>course_id</i>	..	<i>credits</i>
			CS-101		3

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>tot_cred</i>
1	John	Comp. Sci.	100
			103



- Two tables: animals (ID, name); animal_count (animals);
- CREATE TRIGGER increment_animal AFTER INSERT ON animals
FOR EACH ROW
 UPDATE animal_count SET animal_count.animals = animal_count.animals+1;
- CREATE TRIGGER the_mooses_are_loose AFTER INSERT ON animals
FOR EACH ROW
 BEGIN
 IF NEW.name = 'Moose' THEN
 UPDATE animal_count SET
 animal_count.animals = animal_count.animals+100;
 ELSE
 UPDATE animal_count SET
 animal_count.animals = animal_count.animals+1;
 END IF;
 END;
- 若判斷有錯誤發生可寫 SIGNAL SQLSTATE '45000' SET MESSAGE_TEXT ='.....';
- 參考網址：<https://mariadb.com/kb/en/trigger-overview/>



Statement Level Triggers

- Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
 - Use **for each statement** instead of **for each row**
 - Use **referencing old table** or **referencing new table** to refer to temporary tables (called **transition tables**) containing the affected rows
 - Can be more efficient when dealing with SQL statements that update a large number of rows



※ 補充MS-SQL 的Trigger範例

This example creates a trigger that, when an employee job level is inserted or updated, checks that the specified employee job level (**job_lv**) is within the range defined for the job. To get the appropriate range, the **jobs** table must be referenced.

```
CREATE TRIGGER employee_insupd ON employee FOR INSERT, UPDATE AS
/* Get the range of level for this job type from the jobs table. */
DECLARE @min_lv tinyint, @max_lv tinyint, @emp_lv tinyint, @job_id smallint
SELECT @min_lv = min_lv, @max_lv = max_lv, @emp_lv = i.job_lv, @job_id = i.job_id
FROM employee e INNER JOIN inserted i ON e.emp_id = i.emp_id INNER JOIN jobs j ON
j.job_id = i.job_id
IF (@job_id = 1) and (@emp_lv <> 10)
BEGIN
    RAISERROR ('Job id 1 expects the default level of 10.')
    ROLLBACK TRANSACTION
END
ELSE
IF NOT (@emp_lv BETWEEN @min_lv AND @max_lv)
BEGIN
    RAISERROR ('The level for job_id:%d should be between %d and %d.', @job_id, @min_lv,
    @max_lv)
    ROLLBACK TRANSACTION
END
```



When Not To Use Triggers

- Triggers were used earlier for tasks such as
 - Maintaining summary data (e.g., total salary of each department)
 - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
 - Databases today provide built in materialized view facilities to maintain summary data
 - Databases provide built-in support for replication